

PBSMS

Internal Engineering Agent

Engineering Process Specification · Version 2.1 · Model B of two AI specifications · Internal — not a product feature

Item	Value
Document	PBSMS Internal Engineering Agent — Engineering Process Specification
Version	2.1 — Adopted. Closes the pbsms-ai/ artefact gap created by Chapter 47 v2.1.
Status	ADOPTED — APPROVED FOR IMPLEMENTATION. v2.0 adopted 27 August 2026; v2.1 adopted 29 August 2026. Implementation authorisation is staged; see §1.3.
Classification	Internal engineering process. Not an SRS chapter. Creates no system requirement. Tenants are never aware of it.
Requirement prefix	EC- (Engineering Agent). These are process rules binding the team, not software requirements binding the product.
Owner	Whoever holds the Engineering Lead role
Companion	Model A — Tenant AI Assistant, SRS v2.1 Chapter 47, v2.1, adopted 29 August 2026, together with the Domain AI Model Architecture Blueprint v1.0 as its implementation baseline.

Revision Note — v2.0 → v2.1

One change, and it is urgent rather than tidy. Chapter 47 v2.1 and the Architecture Blueprint were adopted on 29 August 2026. The Blueprint establishes a pbsms-ai/ ownership boundary in the repository holding the Model Constitution, the dataset manifests, the training configuration and the blind evaluation corpus. EC-005 and EC-400 were written before any of those existed and did not reach them.

As v2.0 stood, this Agent could modify the blind evaluation corpus — the corpus on which every model release decision under Chapter 47 §47.16 depends. That modification would fail no CI gate, because the gates are computed against that corpus. It would silently invalidate every certification downstream of it, and the failure would be undetectable by exactly the mechanism meant to detect failures.

#	Change	Reason
1	EC-005 replaced — whole-tree write prohibition on pbsms-ai/, replacing the enumerated list of Assistant artefacts	The tree holds both the rules constraining another AI system and the corpora certifying it. An enumerated file list drifts as the tree grows; a boundary does not.
2	EC-400 extended — blind evaluation corpus, dataset manifests and hashes, captured web corpus, adversarial corpus, and the scorers computing the Chapter 47 §47.16.2 metrics	A corpus the Agent can edit is a certification the Agent can forge. Unlike a code defect this one produces no failing test — the test is the thing that changed.
3	EC-501 scope widened to the whole boundary; new EC-507 protected-	EC-501 was already cleared for immediate build. This changes its path list, not its existence.

#	Change	Reason
	boundary check	
4	§1.2, §1.3 and §5 updated to reflect the adopted Chapter 47 artefacts	The tooling table, authorisation table and protected-zone table each predated the model programme.

1. What This Is, and What It Is Not

The Internal Engineering Agent is an AI coding assistant operating inside the PBSMS repository, used by the PBSMS engineering team. It clusters user feedback into a ranked backlog, answers questions about the codebase, detects gaps between the SRS and the implementation, drafts change proposals, and implements approved ones as pull requests on feature branches.

It is not deployed in the product. It has no tenant. It has no user outside the engineering team. It ships nothing by itself.

This document is separate from Chapter 47 for a specific reason: the original proposal treated both as one system, and treating them as one system is the principal risk this separation exists to remove. Model A is a feature a school pays for and uses. Model B is a tool the engineering team uses to build the thing the school pays for. They have different data, different threat models, different approvers, and no shared infrastructure.

1.1 Hard separation from the product

- **EC-001** — The Agent shall not be deployed in, callable from, or reachable by the PBSMS production runtime.
- **EC-002** — The Agent shall hold no credentials to any production database, no production API keys, and no path to any tenant record.
- **EC-003** — The Agent and the Tenant AI Assistant shall share no infrastructure, no credentials and no code path. A compromise of one shall not reach the other.
- **EC-004** — No output of the Agent shall reach a tenant except as reviewed, merged, tested, released code — on exactly the same path as any human engineer's work, with no shortcut and no separate channel.
- **EC-005** — The Agent shall have no write access to the pbsms-ai/ ownership boundary, in whole. This covers the Model Constitution and its compiled artifacts, tool and response schemas, dataset manifests and generators, training configuration, evaluation corpora and scorers, and model release manifests. Read access for question answering under EC-102 is permitted.

The prohibition is a whole-tree boundary rather than an enumerated file list because that tree holds two different things the Agent must not touch: the rules constraining another AI system, and the corpora that certify it. An enumerated list drifts as the tree grows, and the drift is invisible until something has already been edited. v2.0 enumerated, and within two days of Chapter 47 v2.1 the enumeration was already incomplete.

1.2 What this is actually made of

Nothing in this document requires building an AI system. The Agent is Claude Code, operating on the PBSMS repository under the rules that follow. Of the six components below, four are configuration and two are small internal builds. None of them is a model.

Component	What it is	Build or configure
Engineering Agent	Claude Code — run locally by an engineer against the repository, or invoked in GitHub Actions by mentioning @claude on an issue or pull request	Configure
Agent rulebook	CLAUDE.md in the repository root, carrying the prohibitions in §3 and the protected zones in §5. Claude Code reads and follows it on every invocation.	Write once
Backlog	GitHub Issues. Clustered themes are written as issues carrying the affected requirement IDs, tenant count and severity, and are reviewed in the tracker like any other work.	Configure
Protected-zone enforcement	CODEOWNERS plus branch protection rules, backed by the CI checks in §6	Configure
Feedback ingestion and clustering (EC-100, EC-101)	A scheduled job reading tickets, in-app feedback and telemetry; anonymises per EC-300, clusters, writes issues	Build
CI gates (EC-500 to EC-506)	Jobs added to the existing .github/workflows/ci.yml, alongside SAST, the RLS coverage gate and the accessibility gate already there	Build

Two properties of the default behaviour are worth stating, because they mean less has to be engineered than §3 implies:

- Claude Code commits to a branch and returns a link to the pull-request page rather than opening and merging one itself. EC-200's merge boundary is therefore the tool's shipped behaviour, not a constraint bolted onto it. What §3 adds is the reasoning for never changing that, and the branch protection that makes it unchangeable by accident.
- Credentials are repository-scoped — an API key or OAuth token held in repository secrets. This satisfies EC-002 without additional work, on the sole condition that no production credential is ever added to the same secret store. That condition is the one thing in this section a person can quietly break.

What remains genuinely undecided is whether the Agent runs primarily from engineers' machines or primarily in GitHub Actions. Local invocation gives tighter control and no repository-level API credential at all; Actions invocation gives a shared audit trail and lets a backlog issue become a branch without anyone opening a terminal. The rules in this document hold either way, so this can be settled at stage 3 from experience rather than argued now.

1.3 Adoption, amendment and implementation authorisation

This document is adopted as binding engineering process on 27 August 2026. It is not advisory. An engineer working outside these rules is working outside process, in the same sense as bypassing code review.

- Amendment requires a version increment and re-adoption by whoever holds the Engineering Lead role. EC- identifiers are permanent; a withdrawn rule is marked withdrawn and retained.
- EC-005, EC-302 and EC-400 are not amendable by the ordinary process. Changing any of the three requires re-adoption of this document as a whole. They are the three rules whose entire value is that

they cannot be relaxed by whoever is under pressure at the time. v2.1 widened EC-005 and EC-400 rather than narrowing them, which the ordinary process permits.

- Advancing a rollout stage (§11) is an explicit decision by the Engineering Lead against the stated advance condition, recorded. Stages do not advance by drift, by consensus, or because the previous stage has been running a while without incident.

Adoption authorises implementation, staged as follows:

Item	Scope	Authorisation
CLAUDE.md rulebook	The prohibitions in §3 and protected zones in §5, written into the repository root	CLEARED — do this first. It is the cheapest item here and everything else assumes it exists.
CODEOWNERS and branch protection	Two-approver enforcement on protected zones (EC-202)	CLEARED — begin immediately.
EC-501 immutable-suite CI check	Fails any diff modifying or deleting a protected test file, evaluation corpus, dataset manifest or scorer	CLEARED — required before stage 3, the first stage the Agent touches test files.
EC-507 protected-boundary CI check	Fails any Agent pull request whose diff touches pbsms-ai/ at all, including rename or move	CLEARED — build this ahead of everything else here. Chapter 47 §47.18 stage 1 is work inside that tree, and that stage is already cleared to start.
EC-500, EC-502 to EC-506 gates	Remaining CI gates in .github/workflows/ci.yml	CLEARED — begin immediately.
Repository question answering (EC-102)	Engineers using Claude Code against the repo	CLEARED — usable today. It reads code and answers questions; it merges nothing.
Feedback ingestion and clustering (EC-100, EC-101)	The scheduled clustering job	BUILD CLEARED, OPERATION DEFERRED — there are no tenants and therefore no feedback. Build when there is signal to cluster, not before.
Gap detection (EC-107)	SRS-to-implementation comparison	CLEARED — genuinely useful now, with a large SRS and a small implementation.
Stage 3 — test-only pull requests	Added coverage, no production code	GATED on EC-501 existing.
Stage 4 onward — implementation PRs	Code changes on branches	GATED on a human defect-escape baseline existing (§13). That takes roughly a quarter of ordinary development to produce and cannot be shortcut.
Stage 6 — protected zones	Drafting inside the isolation boundary	NOT AUTHORISED by this adoption. A separate decision on its own merits.

Note what is cleared today versus what waits. Every safeguard is cleared immediately; every capability that can cause harm is gated on a safeguard existing first. That ordering is the whole point, and it is also why adoption today costs nothing — there is no tenant feedback to cluster and no baseline to measure against, so the only work actually available right now is the work that makes the rest of it safe.

2. Capabilities

ID	Capability	Output
EC-100	Feedback aggregation	Deduplicated, clustered themes from in-app feedback, support tickets and telemetry friction signals
EC-101	Impact ranking	A ranked backlog with request volume, affected roles, affected tenant count and severity — an input to prioritisation, never prioritisation itself
EC-102	Repository question answering	Answers about the codebase with file and line references, for engineers
EC-103	Change proposal drafting	An RFC: problem, affected SRS requirement IDs, proposed change, schema impact, tenant impact, rollback plan, test plan
EC-104	Implementation	A pull request on a feature branch, with tests and a migration where needed
EC-105	Test authoring	Tests covering its own changes, subject to §5
EC-106	Documentation drift detection	A report of where a merged change has made the SRS or the Frontend Design Specification inaccurate, plus a drafted amendment. It does not amend either document.
EC-107	Gap detection	Requirements in the SRS with no corresponding implementation, no test, or both — reported, not fixed

EC-107 is the "find and fix gaps" capability from the original proposal, split in two. Finding is automated and safe: it is a comparison between two documents the Agent can read. Fixing is not automated, because deciding whether an unimplemented requirement should be implemented, deferred, or deleted from the SRS is a product judgement, and roughly a third of the time the correct fix is to change the requirement rather than the code. Gaps go to the backlog. The backlog goes to a person.

3. The Merge Boundary

- **EC-200** — The Agent shall never merge to main. There is no maturity threshold, no accumulated success rate, and no trusted-mode flag that unlocks this.

The original proposal was that the Agent be trained repeatedly, and once its success rate was high enough, be permitted to push and merge to main. Three specific things are wrong with that, in order of seriousness.

3.1 Self-referential validation proves the wrong thing

A change the Agent wrote, covered by tests the Agent wrote, passing a suite the Agent can influence, demonstrates that the Agent is internally consistent. It demonstrates nothing about whether an RLS policy still isolates one school's records from another, whether a published result is still immutable under FR-RES-030, or whether a fee proration rule still matches a Ghanaian three-term year. Those are questions about intent, and intent is not recoverable from the code that was supposed to express it. This is why §5 makes certain test suites human-authored and immutable to the Agent — they are the only part of the pipeline that is not downstream of the Agent's own understanding.

3.2 Accountability is a legal requirement, not a cultural preference

Ghana's Data Protection Act requires an identifiable controller accountable for decisions about how personal data is processed. PBSMS processes the personal data of children. A merge that changes how that data is stored, exposed, retained or shared is a processing decision and needs a name against it. "The model merged it" is not a position available to a data controller under audit, and the point at which it becomes unavailable is the point at which the incident has already happened.

3.3 Success rate is the wrong instrument

A 99% success rate over a thousand merged changes is ten defects shipped into a shared-database, shared-schema system holding minors' records and school finances. The useful question is not how often the tool is right; it is what happens on the occasion it is wrong. In this architecture the tail risk is cross-tenant data exposure, which is silent, unbounded in scope, not detected by the tenant, and not undone by a rollback — the data has already been seen. A metric that averages that event with a thousand successful button-colour changes is not measuring the thing that matters.

3.4 The rules that follow from this

- **EC-201** — Every Agent pull request requires review and approval by a named human engineer before merge, on the same terms as a human-authored pull request.
- **EC-202** — Pull requests touching a protected zone (§5) require two human approvers, one of whom holds the Engineering Lead role.
- **EC-203** — Deployment to production is human-triggered. The Agent shall hold no deploy permission.
- **EC-204** — The Agent shall not approve any pull request, including one it did not author.
- **EC-205** — The Agent shall not modify CI configuration, branch protection rules, or repository permissions.

4. Feedback Pipeline

Source	Signal	Handling
In-app feedback	Free-text submission from any surface, with role and screen context	Anonymised at ingestion (EC-300)
Support tickets	Tenant-raised issues and requests	Anonymised; tenant identity replaced by an opaque identifier
Telemetry	Abandonment, repeated retries, dead-end navigation, validation-error clustering	Recorded against role and screen, never against an identified user

Source	Signal	Handling
Assistant error reports	Answers reported as incorrect under FR-AIT-602	Question text is personal data — stripped before it reaches the Agent; only the failure class crosses
Tenant success conversations	Notes from calls with schools	Summarised by a human before ingestion, never raw

- **EC-300** — Anonymisation at ingestion. Student names, guardian names, staff names, contact details, admission numbers, invoice numbers, BECE index numbers and any free-text identifier shall be stripped before feedback reaches the Agent's context. The Agent works on the shape of a problem, not on who reported it.
- **EC-301** — Tenant attribution shall be by opaque identifier only. The Agent may know that fourteen distinct tenants reported the same friction. It shall not know which fourteen.
- **EC-302** — No production data in the Agent's context, at any time, in any form — not a sample, not an export, not a temporary copy for debugging. Where a defect cannot be diagnosed without real data, a human diagnoses it.
- **EC-303** — Telemetry shall be behavioural, not personal: events against role and screen, with no user identifier and no tenant identifier beyond the opaque one.

4.1 Cross-tenant contamination

This is the failure mode specific to a shared-schema multi-tenant product, and it is not a security failure — it is a product failure that looks like a feature.

- **EC-310** — Volume is not a mandate. A change requested by many tenants does not thereby become a change every tenant receives. The ranked backlog is an input to a decision made by a named human product owner, who is accountable for that decision.
- **EC-311** — Shared schema means everyone inherits. Any merged change altering a workflow a tenant depends on shall ship behind a feature flag with tenant-level opt-in, defaulting to existing behaviour. The default protects the tenants who did not ask.
- **EC-312** — No mid-term workflow changes. Changes to attendance, assessment, results or invoicing workflows shall not be released to a tenant during an active term without that tenant's explicit consent. A Ghanaian school year has three terms; a workflow change landing in week six of one is disruption, not improvement, however well-requested it was.
- **EC-313** — A minority tenant shall not be outvoted into breakage. Where a highly-requested change would degrade the workflow of tenants who did not request it, the flag default holds them on existing behaviour until they choose otherwise.

5. Protected Zones and Immutable Suites

Some parts of this codebase are load-bearing for correctness properties that cannot be recovered after they break. The Agent may draft changes in most of them, under stricter review. In two of them it may not act at all.

Zone	Why	Rule
RLS policies; any migration	The entire tenant isolation	May draft. Two approvers incl.

Zone	Why	Rule
touching a tenant-owned table	boundary	Engineering Lead. RLS coverage gate and cross-tenant isolation suite must pass.
Tenant context middleware, AsyncLocalStorage scoping, request-scoped DB service	Carries the isolation boundary into every request	Same as above
Authentication, authorisation, Chapter 13 scope resolution	Access control for minors' records	Same as above
Finance ledger, allocation, reversal (Ch 23–25)	Money correctness; maker-checker integrity	May draft. Two approvers, one holding finance domain ownership.
Results immutability and publication gate (FR-RES-020, FR-RES-030)	Academic record integrity; BECE and CSSPS depend on it	Same as above
Subscription metering and billing (TEN-030, TEN-031)	Revenue correctness	Same as above
Chapter 47 Assistant prompts, scope config, evaluation gates	Constraints on another AI system	No access. Prohibited outright (EC-005).
The pbsms-ai/ tree — constitution, schemas, datasets, training config, evaluation corpora, release manifests	Holds both the rules constraining the Domain AI Model and the corpora certifying it	No write access. Prohibited outright (EC-005). Read permitted for EC-102.
Protected test suites and evaluation corpora (below)	The only validation not downstream of the Agent	No access. Prohibited outright (EC-400).

- **EC-400** — The following shall be human-authored and immutable to the Agent: the cross-tenant isolation suite (NFR-QA-020); the finance invariant suite; the results-immutability suite; the Chapter 47 Assistant isolation, grounding and privacy gates; the blind evaluation corpus; dataset manifests and their hashes; the captured web corpus; the adversarial corpus; and the scorers that compute the Chapter 47 §47.16.2 metrics. The Agent may propose additional cases in a pull request that touches nothing else; it may never modify or delete an existing case, corpus, manifest or scorer.

This is the single most important rule in the document. A tool that can edit the tests constraining it is not constrained, and a tool that can edit the corpus certifying a model release can forge that certification. The second case is worse than the first: a modified test suite may still fail somewhere, but a modified evaluation corpus produces a clean green run, because the thing that changed is the measuring instrument. Every other safeguard here degrades gracefully under pressure; this one does not degrade at all, which is why it is an access prohibition enforced by repository permissions and CI rather than a review convention.

- **EC-401** — Agent-authored tests are evidence, not proof. A reviewer shall assume that such a test encodes the Agent's understanding of the change — the same understanding that produced the change. Review shall check the test against the requirement, not against the code it accompanies.

6. CI Gates for Agent Pull Requests

A Agent pull request passes every gate below before a human is asked to look at it. Spending a reviewer's attention on a change that an automated check would have rejected is the fastest way to make reviewers stop reading carefully, and reviewer attention is the actual safety mechanism in this whole design.

ID	Gate	Fails when
—	Existing pipeline	SAST, dependency scan, RLS coverage gate, cross-tenant isolation suite, accessibility gate (pa11y-ci), performance budget — all as already configured
EC-500	Protected-path check	The diff touches a protected zone without the required label and approver count
EC-501	Immutable-suite check	Any protected test file, evaluation corpus, dataset manifest, hash file or scorer is modified or deleted
EC-507	Protected-boundary check	Any file under pbsms-ai/ appears in the diff of an Agent-authored pull request, in any form including rename or move
EC-502	Migration safety check	A destructive migration appears without an explicit, separately reviewed exception
EC-503	Diff ceiling	The change exceeds the reviewable size limit (proposed 400 changed lines) and has not been split
EC-504	Traceability check	The pull request does not reference the SRS requirement IDs it implements or changes
EC-505	Behaviour-change flag check	A user-visible workflow change ships without a feature flag (EC-311)
EC-506	Attribution check	The commit and pull request are not labelled as Agent-authored

7. Change Workflow

1	Signal	feedback, tickets, telemetry, gap detection	AGENT
2	Cluster & rank	deduplicated themes, ranked backlog	AGENT
3	Product decision	what gets built, and for whom	HUMAN
4	RFC draft	problem, impact, schema, rollback, tests	AGENT
5	RFC approval	technical approach signed off	HUMAN
6	Implementation	feature branch, code, tests, migration	AGENT
7	CI gates	§6 – all must pass before review is requested	AUTOMATED
8	Code review	against the requirement, not the code	HUMAN
9	Merge	the accountability record	HUMAN
10	Release	flagged, default off, human-triggered	HUMAN
11	Rollout	tenant opt-in, staged, never mid-term	HUMAN
12	Measurement	adoption, defects, reverts	AUTOMATED
13	Doc drift	SRS / Frontend Spec amendment drafted	AGENT
14	Doc approval	amendment accepted into the baseline	HUMAN

Humans act at seven of the fourteen stages, and at every stage where something becomes irreversible or reaches a tenant. That is the design, not a transitional arrangement pending greater confidence in the tool.

8. Accountability

Role	Accountable for
Engineering Lead	This document; protected-zone approvals; the decision to advance a rollout stage; invoking EC-700
Product Owner	Stage 3 product decisions; which tenants receive which flagged change and when
Reviewing engineer	The correctness of the merged change. Named in the commit trail as author of record for audit and DPA purposes.
Data Protection Officer	That EC-300 anonymisation holds in practice; that no processing decision reaches production without a named human behind it

- **EC-600** — Every merged Agent change shall carry the reviewing engineer's name in the commit trail. That engineer is the author of record.
- **EC-601** — Agent authorship shall be labelled in the pull request and the commit, so the proportion of AI-authored code in the system is measurable rather than folkloric.
- **EC-602** — An incident caused by a Agent-authored change shall be reviewed as a failure of the gates and the process, not of the reviewer alone. A process that depends on a reviewer never having a bad afternoon is not a process.

9. Metrics

Metric	Why it is tracked	What a bad number means
Review rejection rate	Whether review is real	Very low is the warning sign, not the good sign — it suggests rubber-stamping rather than quality
Time-to-review	Reviewer load	Rising means reviewers are saturated and review quality is falling before defect rates show it
Defect escape rate, Agent vs human PRs	The core comparison	Agent exceeding human baseline triggers EC-700
Share of PRs touching protected zones	Blast-radius exposure	Should stay low. Rising means the tool is being pointed at the wrong work.
Reverts within 30 days	Change quality	Rising indicates the gates are passing changes that should not pass
Signal-to-ship latency	The actual benefit being claimed	If this does not fall, the tool is not delivering the thing it was adopted for

- **EC-700** — Suspension condition. If the defect escape rate of Agent-authored merged changes exceeds the human baseline over two consecutive quarters, or if any Agent-authored change causes a cross-tenant data exposure, implementation capability (EC-104, EC-105) shall be suspended immediately and only the analysis capabilities (EC-100 to EC-103, EC-107) retained pending review.

Naming the shutdown condition in advance is what makes it possible to actually invoke. A condition agreed after an incident is a negotiation; a condition agreed before one is a rule.

10. Failure Modes

Failure	How it shows up	Containment
Plausible but wrong implementation	Passes self-authored tests, fails intent	EC-401 review posture; protected suites (EC-400); requirement-referenced review (EC-504)
Reviewer fatigue	Approval rate rises, review time falls	Diff ceiling (EC-503); metrics in §9; this is the hardest failure to detect early and has no good leading indicator
Scope creep in a single PR	A small fix carries an unrelated refactor	Diff ceiling; traceability check; reviewer rejects rather than untangles
Feedback poisoning	A tenant submits feedback crafted to steer the backlog	EC-300 anonymisation; EC-310 human product decision; volume is an input, not an instruction
Protected-zone drift	Zone boundaries erode as the codebase grows	Zone definitions reviewed each quarter by the Engineering Lead; new isolation-critical code added to the list when written, not later
Silent dependency on the tool	Team loses the ability to work without it	Rollout stages are reversible by design; EC-700 suspension must remain survivable
Cross-tenant exposure from a merged change	Possibly never detected by the tenant	The reason for two approvers, immutable isolation suites, and EC-700 naming this as a single-instance trigger

11. Rollout Stages

Stage	Capability enabled	Advance when
1	Feedback clustering and repository Q&A. No pull requests.	The backlog it produces is demonstrably more useful than a product owner reading tickets — which is currently an assumption, not a finding.
2	Gap detection (EC-107), reporting only	Reported gaps are real gaps at an acceptable rate; false-positive rate is tolerable to the team
3	Test-only pull requests — added coverage, no production code touched	Lowest-risk way to calibrate review posture and measure review quality before anything can break
4	Implementation PRs outside protected zones, small diffs, flagged	A human defect-escape baseline exists to compare against (see §13)

Stage	Capability enabled	Advance when
5	Wider implementation, still outside protected zones	Stage 4 defect escape rate at or below human baseline for two quarters
6	Protected zones, drafting only, two approvers	A distinct decision by the Engineering Lead, reviewed on its own merits. Not automatic on the success of stages 1–5.

Stage 3 before stage 4 is deliberate. Test-only pull requests let the team find out how carefully it actually reviews Agent output at a point where a missed defect costs nothing. Discovering the answer to that question at stage 5 is considerably more expensive.

12. On Training, and on Success Rates

Three points, directly on the original proposal, because they change what should be built rather than only what should be permitted.

- You do not train a model on the repository. Retrieval over the current codebase is both cheaper and more correct: a fine-tuned model encodes the repository as it was on the day of training and drifts silently as the code changes, with no signal that it has. Retrieval reads what is there now. The improvement path for this tool is better retrieval, better prompts, better evaluation cases and better gates — not model weights.
- "Success rate" needs a denominator that means something. Measured against self-authored tests, it measures internal consistency. Measured against human review acceptance, it measures usefulness — worth tracking, and it is the review rejection rate in §9. Neither measures safety. Safety is measured only by the suites the Agent cannot touch, which is the entire reason EC-400 exists.
- The merge button is not a bottleneck to be optimised away. It is the accountability record. Removing it does not remove the risk; it removes the name attached to the risk, which under Ghana's Data Protection Act is the one thing that cannot be missing.

13. Open Questions

- Tooling is settled (§1.2) but the invocation model is not — local versus GitHub Actions. Decide at stage 3 from experience, not now from argument.
- EC-400's file-level immutability needs a concrete enforcement mechanism, not just a CLAUDE.md instruction. A rulebook the Agent reads is a rule it can be talked out of; a CI check on the diff and a CODEOWNERS entry are not. Build the check before stage 3, because stage 3 is the first stage where the Agent touches test files at all.
- The 400-line diff ceiling in EC-503 is a guess. Re-set it from stage 3 measurement of how review quality actually degrades with diff size on this team, which is the only place that number can honestly come from.
- No baseline defect escape rate exists for human-authored PBSMS changes, so EC-700's comparison currently has no denominator. Establishing it is a prerequisite for stage 4 and it takes at least a quarter of ordinary development to produce.
- Reviewer fatigue is the real failure mode of this entire design and it has no good leading indicator. The metrics in §9 detect it late. If someone finds a better early signal, that is a more valuable contribution to this document than any other change to it.

- Cost per merged change is unmodelled, so the economic case for adoption is unquantified.
- Whether automated feedback clustering produces a better backlog than a product owner reading tickets directly is an assumption. Stage 1 exists specifically to test it, and stage 1 failing is an acceptable outcome that saves the cost of stages 2 through 6.
- This document, like Chapter 47, has been written but not reviewed by anyone other than its author.

— *End of Document* —